

Using dynamic programming for making decisions about invasive species

Olivier Gimenez^{1*} Abby Keller^{2*} Lucile Marescot^{3*} Carl Boettiger^{2*}
Cassie Speakman^{4,1*}

2026-05-21

¹ CEFE, Univ Montpellier, CNRS, EPHE, IRD, Montpellier, France ² University of California, Berkeley, Berkeley, California, USA ³ CBGP, INRAE, IRD, CIRAD, Institut Agro Montpellier, Univ Montpellier, Montpellier, France ⁴ FRB-CESAB, Montpellier, France

* Corresponding author: olivier.gimenez@cefe.cnrs.fr

Invasive alien species (IAS) pose significant ecological, economic, and public health challenges. Robust methods are needed for guiding effective management. Overall, we hope to demonstrate the flexibility and utility of structured decision making in invasive species management [Abby: I think of structured decision making as being more about *process* (i.e., PrOACT steps of problem framing, eliciting objectives and alternatives from decision makers, etc), whereas I think what we've done here is given an overview of decision theoretic methods that are useful in one step of the structured decision making process? I think Hemming et al. make some useful definition distinctions: <https://doi.org/10.1111/cobi.13868>]. [Cassie: Agreed; here we're presenting approaches that can make the solving of complex decision problems as part of the SDM process. Could it instead be reframed as 'Overall, we hope to demonstrate the flexibility and utility of dynamic programming for solving the complex decisions inherent to invasive species management.' ?]

²² Target journal: Journal of Animal Ecology

²³ *Keywords:* Adaptive management, Decision theory, Invasive alien species, Stochastic dynamic
²⁴ programming, Structured decision making, Uncertainty

25 Introduction

26 Invasive species are a major global concern, with wide-ranging impacts on ecosystems,
27 economies, and public health (? , ?). Their management is inherently challenging, as it requires
28 repeated decisions under uncertainty, limited budgets, and imperfect knowledge of ecological
29 dynamics, which are often characterized by rapid spread, strong population fluctuations, and
30 broad spatial distributions. In practice, managers must decide not only whether to intervene,
31 but also when, where, and how intensively, while balancing short-term operational costs
32 against long-term socio-ecological benefits (?).

33 This makes invasive species management a fundamentally dynamic decision problem. Predict-
34 ing invasion dynamics or mapping risks is not sufficient unless these predictions are explicitly
35 linked to management objectives, feasible actions, and measurable consequences. Decision-
36 making is further complicated by the need to account for multiple, sometimes competing,
37 criteria, including ecological effectiveness, economic costs, operational feasibility, and potential
38 unintended consequences such as non-target impacts. In addition, management actions are em-
39 bedded in social contexts, where stakeholder perceptions, acceptability of control measures, and
40 underlying environmental values can strongly influence both implementation and outcomes (? ,
41 ? , ?)

42 Structured decision-making (SDM) and adaptive management (AM) have long been advocated
43 as frameworks to address the complexity of ecological management problems (? , ?). These
44 approaches emphasize the explicit formulation of objectives, the evaluation of alternative
45 actions, and the incorporation of uncertainty into decision processes. However, despite their
46 conceptual appeal, they remain underused in practice, in part because of the technical challenge
47 of formally identifying, quantifying, and optimizing trade-offs among competing objectives
48 over time. Stochastic dynamic programming (SDP), a core tool of decision theory, provides a
49 powerful framework to address this challenge by treating management as a sequential decision

problem under uncertainty. In this framework, optimal actions are derived as a function of the current state of the system (e.g. infestation size, spatial extent, or local density), while explicitly accounting for stochastic dynamics and the future consequences of present decisions. SDP naturally accommodates a wide range of management actions, including prevention, surveillance, containment, control, eradication, or inaction, and allows their outcomes to be evaluated in terms of long-term ecological, economic, or social objectives (?).

SDP has been successfully applied to a variety of invasive species management problems, including optimizing prevention versus control strategies for zebra mussels (?), designing control policies for invasive plants such as leafy spurge (?), identifying state-dependent strategies for spatially structured invasions (?), optimizing search effort for invasive insects (?), and comparing alternative strategies across invasion stages (?). More recent developments, including partially observable Markov decision processes (POMDPs), allow decision-making under imperfect detection, while connections with reinforcement learning highlight promising avenues for scaling these approaches to complex, data-rich systems (e.g. Haubrock et al. (?)) [Abby: Is this paper an example of reinforcement learning for invasive species management?]. [Cassie: For POMDP, the Waring et al 2023 MEE paper could be a good one to point to]

Despite these advances, the practical implementation of SDP and related approaches remains limited in applied animal ecology. This gap is partly due to the perceived complexity of these methods and the lack of accessible, step-by-step guidance for translating ecological problems into operational decision models.

In this paper, we aim to bridge this gap by providing a practical guide to the use of stochastic dynamic programming and its extensions for invasive species management. Using the coypu (*Myocastor coypus*) as a running example, we illustrate how to formulate decision problems, implement models, and derive optimal management strategies across a range of ecological contexts, from single populations to metapopulations. We first introduce the core concepts of structured

75 decision making and stochastic dynamic programming through a minimal, fully transparent
76 example. We then develop a series of worked tutorials of increasing complexity, including
77 spatial dynamics, imperfect detection using partially observable Markov decision processes
78 (POMDPs), and model uncertainty addressed through adaptive management. Throughout, we
79 provide reproducible implementations in R to facilitate uptake and adaptation by practitioners.
80 [Cassie: I feel like we should keep using the acronymns here (i.e. SDP, POMDP) though if we
81 use SDM too it might become confusing..]

82 [Cassie: Given that we have the sections below to detail SDM, SDP and AM in more detail,
83 should we maybe focus the intro a little more on why we would want to manage invasive
84 species, challenges in making the management decisions and just briefly introduce the SDM,
85 SDP and AM? I haven't read the 'How to' style papers for JAE so not sure how much redundancy
86 they usually have]

87 **Explanation of the method**

88 "A very brief historical background. Why is the method useful for investigating or analysing
89 the topic in question? Provide some details on the method's applications."

90 We do not want to rewrite Lucile's paper...

91 **Structured decision making**

92 [Abby: Maybe a good place to describe structured decision making and how the methods we
93 present are mostly optimization methods useful for evaluating the consequences of actions (i.e.,
94 we're not addressing the aspects of SDM such as problem framing, eliciting objectives; we're
95 instead assuming this has already been done)] [Cassie: I've briefly tried to address these points
96 below; references need work]

97 (?, ?)

Structured decision making (SDM) provides a framework for decision makers to approach complex decision problems by clearly establishing the decision context, formulating clear and measurable objectives, identifying the available management alternatives (actions), identifying the key uncertainties that might impact the management outcomes, and managing trade-offs and constraints (Gregory et al. 2012). The SDM framework thus breaks complex decisions into smaller components, producing decisions that are more transparent and robust than decisions made without a structured approach.

In this paper, we present some of the methods that can be used to solve complex decision problems formulated using an SDM process by evaluating the consequences of alternative management actions in the context of the objectives, uncertainty and trade-offs. Naturally, that means we are assuming that the decision context has already been defined, the objectives elicited, and alternative actions identified, before implementing the below approaches. [Cassie: Here could be a good point to direct people to the Hemming et al 2021 paper]

Stochastic dynamic programming (Olivier)

Here we present a minimal stochastic dynamic programming (SDP) example. The aim is purely pedagogical: to show how states, actions, transition probabilities, rewards, and time horizon combine to produce an optimal decision rule.

We consider a stylized invasion management problem with three qualitative states describing the invasion level:

- Eradicated: no individuals detected (possible re-invasion risk)
- Contained: invasion present but at low abundance
- Established: invasion widespread with high impacts

At each time step, the manager chooses between two actions:

- Monitor: surveillance only (low cost, little effect)




Figure 1: Schematic representation of the Bellman recursion in stochastic dynamic programming. The value of being in a given state at time t , $V(X_t)$, is defined as the maximum, over all possible actions, of two components: (i) the immediate reward obtained by taking action A_t in state X_t , and (ii) the expected discounted future value of subsequent states. The latter corresponds to the average (weighted by transition probabilities) of the optimal future values $V(X_{t+1})$ that can be reached from the current state-action pair, scaled by a discount factor $\beta < 1$ to reflect that future payoffs are less valuable than present ones. This recursive equation highlights the fact that the value of a decision is not only determined by its immediate payoff, but also by how it shapes future opportunities. The optimal action is therefore the one that maximizes the sum of current rewards and expected future benefits.

- Control: active management campaign (higher cost, improves state)

State dynamics are described by action-dependent transition matrices

$$P_a(s, s') = \Pr(S_{t+1} = s' \mid S_t = s, A_t = a).$$

[Abby: Highlighting that we likely want to make a decision about shared syntax. For example, we use x and s interchangeably for the state]

Transition matrices describe how the invasion state evolves probabilistically from one time step to the next under each management action. Each row corresponds to the current state and each column to the next state, with entries giving transition probabilities. We build a transition matrix for each action.

Under monitoring, the transition is:

		Erad.	Cont.	Estab.
$P_{\text{Monitor}} =$	Erad.	0.95	0.05	0.00
	Cont.	0.00	0.80	0.20
	Estab.	0.00	0.00	1.00

[Mettre ce tableau avant l'équation, pour insister d'abord sur l'intuition.]

If the system is eradicated, it remains so with probability 0.95 but may revert to a contained state with probability 0.05, reflecting a small risk of re-invasion. When the invasion is contained, it persists with probability 0.80 but may deteriorate to an established state with probability 0.20, indicating possible population growth without intervention. Once established, the system remains established with probability 1, meaning monitoring alone cannot reverse the invasion.

In contrast, control shifts the system toward better states:

		Erad.	Cont.	Estab.
$P_{\text{Control}} =$	Erad.	0.98	0.02	0.00
	Cont.	0.20	0.75	0.05
	Estab.	0.05	0.75	0.20

From an eradicated state, the system remains eradicated with probability 0.98 and only rarely becomes contained (0.02). When contained, control can eradicate the population (0.20), maintain it at low levels (0.75), or fail and allow establishment (0.05). Finally, when the invasion is established, control improves the situation in most cases, leading to a contained state with probability 0.75, occasionally achieving eradication (0.05), but sometimes failing to reduce impacts (remaining established with probability 0.20).

Together, these transition matrices encode both ecological uncertainty and the imperfect effectiveness of management actions.

[Mettre l'équation ici.]

Next we define rewards. Rewards are defined as the negative of total costs, so that maximizing expected reward is equivalent to minimizing management and damage costs.

$$R(s, a) = -(\text{damage}(s) + \text{cost}(a)).$$

Damage depends on the invasion state (higher when the invasion is established), while management cost depends on the chosen action (higher for active control than monitoring). The reward matrix is constructed by combining these two components for all state-action pairs, yielding the immediate payoff used in the dynamic programming recursion.

		Monitor	Control
$R =$	Erad.	-1	-6
	Cont.	-5	-10
	Estab.	-15	-20

153 (Rows correspond to the current invasion state and columns to the management action.)

154 For example, when the invasion is eradicated, monitoring yields a reward of -1, reflecting
155 only minor surveillance costs, whereas applying control in the same state yields -6 due to
156 higher intervention costs. When the population is contained, rewards are -5 for monitoring
157 and -10 for control, representing moderate damages combined with the respective management
158 costs. Finally, when the invasion is established, rewards drop to -15 under monitoring and
159 -20 under control, reflecting both high ecological impacts and, in the latter case, substantial
160 intervention effort. Overall, this reward structure captures the short-term trade-off between
161 ecological impacts and management effort that underlies the decision problem.

162 Solving the finite-horizon decision problem [ça va un peu vite, et il faut introduire la figure]
163 yields the optimal value function $V_t(s)$, which represents the expected cumulative reward (i.e.,
164 negative total cost) obtained when starting from state s at time t and following the optimal
165 policy thereafter. Values are more negative when the invasion is established, reflecting higher
166 expected damages. Moreover, values increase (become less negative) as the horizon shortens
167 because fewer future costs can accumulate. [Détailier le calcul?]

		$t = 1$	$t = 2$	$t = 3$	$t = 4$	$t = 5$	$t = 6$
$V_t(s) =$	$s = \text{Erad.}$	-7.60	-5.50	-3.69	-2.20	-1	0
	$s = \text{Cont.}$	-33.78	-27.43	-19.96	-12.0	-5	0
	$s = \text{Estab.}$	-49.26	-42.05	-34.47	-26.8	-15	0

168 The associated optimal policy $\pi_t(s)$ specifies the best action to take in each state s at each
 169 time step t . Overall, the policy indicates that monitoring is typically optimal when the system
 170 is eradicated, whereas control is recommended when the invasion is established due to the
 171 high expected damages. In the contained state, the optimal decision varies with the time
 172 horizon, reflecting the trade-off between immediate intervention costs and the risk of future
 173 deterioration.

		$t = 1$	$t = 2$	$t = 3$	$t = 4$	$t = 5$
$\pi_t(s) =$	$s = \text{Erad.}$	Monitor	Monitor	Monitor	Monitor	Monitor
	$s = \text{Cont.}$	Control	Control	Monitor	Monitor	Monitor
	$s = \text{Estab.}$	Control	Control	Control	Control	Monitor

174 This toy example highlights key features of SDP:

- 175 • **State dependence:** optimal decisions depend on invasion level.
- 176 • **Trade-off:** control is costly but reduces future risk.
- 177 • **Time horizon effects:** as the horizon shortens, actions may become more short-term
 178 oriented.
- 179 • **Uncertainty:** decisions account for probabilistic transitions, not deterministic outcomes.

180 [Rappeller les différentes étapes, reprendre le papier de Lulu.]

181 Things to consider for before using this method

182 Applying stochastic dynamic programming (SDP) to ecological problems requires careful for-
 183 mulation of the decision context, as well as practical considerations regarding model structure,
 184 data, and computation.

Problem formulation

A prerequisite for SDP is a clear definition of the decision problem. This includes specifying management objectives, identifying feasible actions, defining the state variables that describe the system, and constructing a reward function that captures the trade-offs among ecological, economic, and social outcomes. Because SDP optimizes decisions given these elements, poorly specified objectives or unrealistic action sets can lead to misleading recommendations.

State definition and discretization.

SDP typically requires a finite state space, whereas ecological systems are often described by continuous variables (e.g. abundance, spatial extent). Discretization is therefore necessary, but introduces a trade-off between realism and tractability. Fine discretization improves accuracy but increases computational cost, while coarse discretization may obscure important system dynamics. Careful sensitivity analyses are often needed to ensure robust conclusions.

Model specification and uncertainty.

The performance of SDP depends on the specification of transition dynamics and reward functions, which are often uncertain in ecological systems. These uncertainties may arise from limited data, structural assumptions, or imperfect knowledge of management effectiveness. Ignoring uncertainty can result in suboptimal policies, motivating extensions such as adaptive management or partially observable Markov decision processes (POMDPs), which explicitly account for learning and imperfect information.

Computational constraints.

A major limitation of SDP is the “curse of dimensionality”, whereby the size of the state and action space grows rapidly with system complexity. This can restrict the number of state variables, spatial resolution, or management options that can be included in a model. In practice,

this often requires simplifying assumptions, state aggregation, or the use of simulation-based approaches to approximate optimal policies. [Cassie: Discuss the heuristic sampling (HS) method that approximates the optimal policy for any starting state from Nicol Chades 2011?]

Data requirements and interpretation.

Although SDP can be implemented using simple or expert-based models, its relevance depends on the availability and quality of data describing system dynamics, management effects, and costs. In data-limited contexts, SDP may still provide valuable qualitative insights into decision structure and trade-offs, but quantitative recommendations should be interpreted with caution. With large state spaces, and particularly for POMDPs, model outputs can become unmanageable and uninterpretable, making it necessary to reduce the state (and observation) space to the minimum necessary to capture the state dynamics. Alternatively, Ferrer-Mestres et al. (2020) leverages AI to present an approach that starts with all states of an SDP and iteratively aggregates states to a smaller number while minimising the performance impacts.

[Cassie: could we split data reqs and interpretation? We could then discuss a little the process of getting the data for objectives and management alternatives (i.e. it's mostly through expert elicitation not from scraping data from papers)]

Worked examples

Here explain that we use coypr IAS in Europe as a “fil rouge”.

Single population (Olivier)

The objective of this example is to determine an optimal regulation policy for coypr based on their total abundance. We consider a stylized management problem in which a decision maker chooses, at each time step, the intensity of a control effort $u \in [0, 1]$. Low effort corresponds

230 to minimal intervention, while high effort represents intensive removal. The management
 231 objective is to balance three competing components: ecological damages that increase with
 232 population size (e.g. crop losses, dike weakening, nuisance), management costs that increase
 233 with effort, and an additional penalty if abundance exceeds a tolerable level.

234 In contrast to the previous example with discrete invasion states, the system is here described
 235 by a continuous state variable, the population size N . Because dynamic programming requires
 236 a finite state space, abundance is discretized onto a grid ranging from 0 to the carrying capacity
 237 K . In the implementation, we set $K = 2000$ individuals and use a grid step of 20 individuals.
 238 Control effort is also discretized, from 0 to 1 in increments of 0.05.

239 Population dynamics follow a density-dependent growth model with removals due to control,

$$N_{t+1} = N_t + rN_t \left(1 - \frac{N_t}{K}\right) - q(u)N_t,$$

240 where the first term describes logistic growth and $q(u)$ is the fraction of individuals removed as
 241 a function of effort. We use a somehow saturating removal function,

$$q(u) = 1 - \exp(-\alpha u),$$

242 with $\alpha = 1.5$, which implies diminishing returns of control effort (Figure ??). For example, an
 243 effort of $u = 0.5$ removes about 53% of the population, whereas maximum effort $u = 1$ removes
 244 about 78%, reflecting the fact that complete eradication is rarely achievable in a single step.

245 Rewards are defined as the negative of total costs, so that maximizing expected reward is
 246 equivalent to minimizing long-term damages and management expenses (Figure ??):

$$R(N, u) = -(\text{damage}(N) + \text{cost}(u) + \text{penalty}(N)).$$




Figure 2: General title. (a). Removal fraction $q(u)$ as a function of control effort u . (b) Immediate cost components as functions of abundance. Damage increases nonlinearly with population size, leading to rapidly rising total costs (damage + penalty). The dotted line marks the tolerable threshold ($N_{tol} = 400$) above which an additional penalty is applied in the decision model, although it is not shown here for clarity. Control cost is illustrated for an example effort level ($u = 0.6$). (c) Optimal policy $u^*(N)$ starting from $N_0 = 800$. (d) An example trajectory under the optimal policy, faceted into abundance and effort, starting from $N_0 = 800$.

247 Damages increase non-linearly with abundance, here using a simple quadratic form
 248 $\text{damage}(N) = 0.1N + 0.001N^2$, reflecting the common ecological pattern whereby impacts
 249 accelerate at high density. Management costs increase with effort and are assumed convex
 250 ($\text{cost}(u) = 50u + 200u^2$), capturing increasing marginal costs of intensive interventions. Finally,
 251 a penalty is applied when abundance exceeds a tolerable threshold $N_{\text{tol}} = 400$, representing
 252 situations where ecological or social impacts become unacceptable. The penalty increases
 253 linearly beyond this threshold according to

$$\text{penalty}(N) = \begin{cases} 0, & N \leq N_{\text{tol}}, \\ 2000 \frac{N - N_{\text{tol}}}{N_{\text{tol}}}, & N > N_{\text{tol}}. \end{cases}$$

254 This formulation implies that the penalty is zero below the threshold and then rises proportion-
 255 ally to the exceedance. For example, an abundance of $N = 800$ yields a penalty of 2000, thereby
 256 strongly incentivizing control once populations grow beyond acceptable levels.

257 The decision problem is solved using infinite-horizon discounted value iteration (with discount
 258 factor $\gamma = 0.99$). The resulting value function represents the expected long-term net benefit of
 259 starting from each abundance level when optimal decisions are applied thereafter. The optimal
 260 policy $u(N)$ typically exhibits a threshold-like pattern: low effort at low abundance where
 261 damages are limited, progressively stronger intervention as abundance increases, and intensive
 262 control once the population approaches or exceeds the tolerable level. This behavior reflects a
 263 general ecological intuition — management intensity should scale with both population size
 264 and risk — and illustrates how stochastic dynamic programming formalizes this principle
 265 within a transparent decision framework.

266 The decision problem is solved using infinite-horizon discounted value iteration (discount
 267 factor $\gamma = 0.99$). This yields (i) an optimal stationary policy, mapping each abundance class
 268 to a recommended control effort, and (ii) a value function (not shown) that summarizes the

long-term expected net benefit when the policy is followed.

Because damages and the penalty increase sharply with abundance, the optimal rule is “threshold-like” (top panel in Figure ??). Effort is low at small N , then increases rapidly as N approaches the tolerable threshold $N_{\text{tol}} = 400$, and remains high for large N . Ecologically, this tells us that when the population is small, impacts are limited and additional control has diminishing returns, whereas once abundance enters an unacceptable zone, strong intervention is warranted to prevent sustained high impacts.

The bottom panel of Figure ?? illustrates the dynamical consequences of applying this rule to a trajectory starting from $N_0 = 800$. Control effort $u(t)$ adjusts immediately to the current abundance class, and abundance $N(t)$ responds through the logistic dynamics with removals. In this deterministic grid-based version, the system typically moves quickly toward a lower-abundance regime and then stabilizes around a band where the marginal benefit of additional control balances its marginal cost.

Metapopulation (Olivier)

We extend the single-population coy-pu example to a spatially structured invasion in which several sites are connected through dispersal. The aim is to explore how management effort should vary through time when local demographic processes and spatial spread interact. We consider a network of several sites that can be interpreted as cities or management units. At each monthly time step, a decision maker chooses a global management action corresponding to a harvest intensity applied simultaneously across all sites. Local populations then grow according to density dependence, are partially removed by management, and may subsequently undergo stochastic extinction or colonization depending on local abundance and the occupancy of neighboring sites.

At time t , each site i is either occupied or empty, represented by the binary variable $X_i(t) \in$

293 $\{0, 1\}$. When a site is empty its abundance is zero, whereas when it is occupied its population
 294 size follows discrete logistic growth with intrinsic rate r and carrying capacity K . Denoting
 295 abundance by $N_i(t)$, post-growth abundance is

$$N_i^{\text{grow}}(t+1) = N_i(t) + r N_i(t) \left(1 - \frac{N_i(t)}{K}\right).$$

296 Management then removes a fraction h_a of the population, so that

$$N_i(t+1) = N_i^{\text{grow}}(t+1) (1 - h_a).$$

297 This formulation captures the idea that control reduces numbers but does not eliminate under-
 298 lying demographic regulation.

299 Occupancy is updated stochastically based on local abundance and spatial context (Figure ??).
 300 When a site is occupied, extinction probability decreases with abundance according to

$$p_{\text{ext}}(N_i) = \exp(-\alpha N_i),$$

301 so that small populations are more vulnerable to disappearance. When a site is empty, coloniza-
 302 tion probability increases with the number of occupied neighbors

$$k_i(t) = \sum_{j \neq i} A_{ij} X_j(t),$$

303 through

$$p_{\text{col}}(k_i) = 1 - (1 - \beta)^{k_i}.$$

Together, these rules link local demography with spatial dynamics and create feedbacks between sites.




Figure 3: General title. (a) Extinction $p_{ext}(N)$. (b) Colonization $p_{col}(k)$. (c) Cumulative cost: constant vs SDP-optimal strategies. Representation of the cost components on the log scale. Damage costs increase linearly with total abundance across sites, whereas control costs are fixed action-specific intercepts (none/moderate/strong). Stronger control increases immediate spending but can reduce future abundance and damages. (c) Optimal global action as a function of time and invasion level (measured as the number of occupied sites). Colors indicate the action recommended by the finite-horizon SDP (None/Moderate/Strong). This visualization highlights how stronger control becomes optimal as the invasion occupies more sites, and how recommendations may soften near the end of the planning horizon. (d) Cumulative costs over time under constant strategies (none, moderate, strong) versus the SDP-optimal policy. Costs combine damages proportional to total abundance and an action-specific control cost. (e) Site-level abundance dynamics under constant strategies and the SDP-optimal policy. Each panel corresponds to a site; trajectories highlight how spatial heterogeneity and network-driven recolonization can create uneven benefits of control across locations. (e) Containment success (sites with ≤ 5 coyus).

Because the transition probabilities induced by these ecological rules are analytically intractable, we approximate them using repeated simulations. Starting from a given state and management action, we simulate the ecological dynamics many times and record the frequency of each possible next state. In practice, the estimated probability of transitioning from state s to state s'

310 under action a is the proportion of simulations that lead to s' :

311 This Monte Carlo approach allows the decision model to incorporate nonlinear and stochastic
312 ecological processes without requiring an explicit analytical expression for transition probabili-
313 ties.

314 Management decisions are evaluated through a cost function that balances ecological damages
315 and intervention effort (Figure ??). At each time step, total cost is defined as

$$C(s, a) = c_d \sum_{i=1}^n N_i(t) + c_a(a),$$

316 where c_d represents the damage cost per individual (here $c_d = 1$) and $c_a(a)$ is the management
317 cost associated with action a . In this example, we consider three possible actions corresponding
318 to no control, moderate control, and strong control, with costs

$$c_a(a) \in \{0, 5, 10\}.$$

319 The reward used in the dynamic programming recursion is defined as the negative of this cost,

$$R(s, a) = -C(s, a),$$

320 so that maximizing expected reward is equivalent to minimizing long-term total costs.

321 The resulting policy is illustrated in Figure ?. When few sites are occupied and populations
322 remain small, limited intervention is often sufficient because damages are modest and extinction
323 may occur naturally. As more sites become occupied or local abundances increase, stronger
324 control becomes optimal in order to reduce future spread and prevent escalating impacts.
325 Temporal variation in recommended actions arises because the finite planning horizon reduces

the benefit of aggressive intervention late in the time period.

Comparing forward simulations under the optimal policy with simulations under constant strategies highlights the value of adaptive management (Figure ??). The optimal strategy typically reduces cumulative costs by applying effort strategically in time rather than maintaining a fixed level of control.

Spatial trajectories further illustrate how coordinated management alters invasion dynamics across the network, with reductions in abundance and occupancy emerging progressively across sites (Figure ??).

In perspective, say that this can be used to model competition or predator-prey, or host-pathogen / disease thing. See SDP-species-interactions/ in Olivier biblio repo.

Accounting for observation error with POMDP (Cassie)

In an ideal world, invasive species management would be based on perfect knowledge of the state of the system. Yet ecological systems are inherently difficult to monitor, leading to imperfect knowledge of the system and its state. While ongoing monitoring can reduce uncertainties associated with the state of the system, invasive species management often requires timely decisions. In these cases, partially observable Markov decision process (POMDP) models can be used. POMDPs build on MDPs, dealing with uncertainty in both the model and state of the system. However, POMDPs are more computationally and mathematically challenging to solve due to the continuous and high-dimensional state space. A detailed primer on developing POMDPs for ecological problems already exists (Chades et al 2021), so here we will keep our description simple.

POMDPs differ from MDPs by including an observation state (O), observation function (Z), and initial belief state (b_0). The observation state includes the set of observations perceived by the manager, while the observation function represents the conditional probability of a manager

observing o' given that action a lead to state s' . In cases where state uncertainty cannot be reduced, the observation function is independent of the actions taken, and thus dependent only on state s' . The initial belief represents the believed state of the system at a given time.

Because the states are themselves unobservable, the state of the system must be tracked with belief states. At any point in time, the actual (unobservable) state of the system influences the immediate returns, state transitions, and observations, but do not influence actions. The belief state is updated based on observations, which subsequently inform the action(s) to be selected. Lastly, actions drive the state transitions, rewards, and may influence the observations. Some may find it helpful to think of POMDPs as an optimal control framework for hidden Markov models (HMMs), where the true system state is latent and must be inferred from observation states (Chades et al 2021).

[I think having a diagram to show the difference between MDP and POMDP, like Figure 1 in Chades et al 2021 or Fig 2-3 in Williams & Brown 2022, could be helpful]

Chades et al (2021) identified three types of management problems that POMDPs can be used to solve:

- Exploiting current knowledge or reducing uncertainty: Classic dilemma where managers have to choose between investing in reducing state uncertainty or exploiting current knowledge.
- Managing under imperfect detection: Assumes that managers have no means of reducing uncertainty, but managers can invest in management actions for which effectiveness varies across the state space.
- Adaptive management: Where the dynamics of the system are not observable.

The authors included an example for the first type of problem from a conservation perspective. Here, we'll focus on an example for managing under imperfect detection, continuing with our coypu case study.

375 Going back to a simple example in invasive species management, imagine we have . . . [Occu-
376 pancy model with uncertain detectability]

377 The invasive species can either be established or absent, with possible observations defined as
378 present or absent denoting if the species was observed in a given time step.

379 POMDPs extend the standard Markov decision process framework to situations in which the
380 true state of the system is not known with certainty. This is particularly relevant for invasive
381 species management, where managers rarely know the exact status of an invasion. Absence of
382 observations does not necessarily mean true absence, and management decisions often have to
383 be made from imperfect monitoring data.

384 Here, we illustrate the approach with a deliberately simple toy example inspired by coypru
385 (*Myocastor coypus*) management. The system can occupy one of three ecological states:

$$s_t \in \{\text{Eradicated}, \text{Contained}, \text{Established}\}.$$

386 These states represent a gradient from no animals present, or at least no animals remaining
387 after management, to a low-abundance contained population, and finally to an established
388 population generating high ecological or socio-economic damage. At each time step, the
389 manager chooses one of three actions:

$$a_t \in \{\text{Do nothing}, \text{Fertility control}, \text{Lethal control}\}.$$

390 Each action affects the probability that the system moves from its current state to a future state.
391 This is described by a transition matrix:

$$P_a(s, s') = \Pr(s_{t+1} = s' \mid s_t = s, a_t = a),$$

where rows correspond to current states and columns to next states. For example, doing nothing is cheap but allows contained populations to become established, while an established population remains established in the absence of control. Fertility control partially reduces invasion pressure, whereas lethal control gives the highest probability of moving the system back towards the eradicated or contained states.

Unlike a fully observable MDP, the manager does not directly observe the state s_t . Instead, they receive an observation:

$$o_t \in \{\text{Absent}, \text{Present}\}.$$

The observation process is governed by detection probabilities:

$$O_a(s, o) = \Pr(o_t = o \mid s_t = s, a_t = a).$$

In the example, detection depends both on the true state and on the action taken. Detection is lower when no management is implemented, reflecting reduced monitoring effort. In contrast, both fertility control and lethal control are assumed to involve field activities that increase the probability of detecting coypu. Detection also increases with invasion status: observations of presence are more likely when the population is established than when it is contained.

Because the true state is uncertain, decisions are based on a **belief state**, denoted b_t . This belief state is a probability distribution over the possible ecological states:

$$b_t(s) = \Pr(s_t = s \mid \text{past actions and observations}),$$

with

$$\sum_s b_t(s) = 1.$$

408 For instance, a belief state such as

$$b_t = (0.10, 0.15, 0.75)$$

409 means that, given all available information, the manager assigns probabilities 0.10, 0.15 and 0.75
 410 to the system being eradicated, contained and established, respectively. In ecological terms, the
 411 manager strongly suspects that the invasion is established, but still acknowledges uncertainty.
 412 The objective is to choose actions that maximise expected long-term reward. Rewards combine
 413 ecological or socio-economic damage and management cost. In this example, the reward is
 414 defined as a negative cost:

$$R(s, a) = -\{D(s) + C(a)\},$$

415 where $D(s)$ is the damage associated with state s , and $C(a)$ is the cost of action a . Thus,
 416 higher rewards correspond to better outcomes, because they represent lower total damage and
 417 lower management cost. Eradication has no damage, containment has moderate damage, and
 418 establishment has high damage. Doing nothing has no management cost, fertility control has
 419 an intermediate cost, and lethal control has the highest cost.

420 The POMDP solution identifies an optimal policy, that is, a rule mapping belief states to actions:

$$\pi^*(b_t) = \arg \max_a V(b_t, a),$$

421 where the value of a belief state is the maximum expected discounted reward over current and

422 future decisions:

$$V(b_t) = \max_{a_t} \left[\sum_s b_t(s) R(s, a_t) + \gamma \sum_o \Pr(o_{t+1} = o \mid b_t, a_t) V(b_{t+1}) \right].$$

423 Here, γ is the discount factor, set to 0.95 in the example, meaning that future outcomes are
424 valued almost as much as immediate outcomes. The updated belief state b_{t+1} is obtained by
425 combining the transition model and the observation model using Bayes' rule. In words, the
426 manager first predicts how the system may change under the chosen action, then updates their
427 belief after observing whether coypru are detected or not.

428 The model can be solved using standard POMDP algorithms. In the code, the same problem is
429 implemented with the `pomdp` package and with `sarsop`. The solution is represented by a set of
430 α -vectors. For any belief state b , the value of each candidate policy vector is obtained by a dot
431 product: [Cassie: I think it would be better to place the description of SARSOPs and describe
432 the alternative solvers here rather than later on in the active AM section]

$$V_i(b) = b^\top \alpha_i,$$

433 and the optimal action is the one associated with the vector giving the highest value:

$$\pi^*(b) = a_{i^*}, \quad i^* = \arg \max_i b^\top \alpha_i.$$

434 This makes the example useful pedagogically: rather than assuming that the invasion status is
435 known, the decision explicitly depends on the manager's uncertainty. When the belief state
436 places high probability on establishment, more intensive control may become optimal despite its
437 higher immediate cost. Conversely, if the system is probably eradicated or only weakly invaded,
438 the optimal policy may favour less costly actions. The POMDP therefore formalises a key

challenge in invasive species management: acting adaptively when both ecological dynamics and monitoring information are uncertain.

Reducing model uncertainty w/ adaptive management (Abby)

Decision theoretic methods can be used to make control decisions not only in the face of uncertainty, but also to reduce uncertainty over time. Adaptive management, or “learning by doing”, involves making decisions to achieve a management objective while simultaneously learning about system dynamics to improve future management success (?, ?, ?). Adaptive management approaches differ based on whether they learn only from the past (passive adaptive management) or anticipate future learning (active adaptive management) (?).

Here we apply adaptive management strategies to inform control decisions for a single coyote population and to learn about and reduce uncertainty in the per-capita growth rate. We will use the following discrete-time population model, where r is the intrinsic rate of growth; K is the carrying capacity; $q(u)$ is the harvest rate as a function of action u ; R is the per-capita growth rate; m describes whether the per-capita growth rate is linear, concave, or convex; and ϵ is log-normally distributed noise:

$$N_{t+1} = N_t(1 + R(N_t)) - q(u) \times N_t + \epsilon$$

$$R(N_t) = r \times (1 - (N_t/K)^m)$$

If $m > 1$, most density—dependent change occurs at high population levels (close to the carrying capacity), which is theoretically common for species with life history strategies typical of large mammals. If $m < 1$, most density-dependent change occurs at low population levels, which is theoretically common for species with life history strategies typical of insects and some fishes (?).

460 We will consider a model set with five possible values of m . This model bounds, but does not
 461 contain, the true model we will use to simulate dynamics in the passive adaptive management
 462 approach (Figure 1A). These different values of m result in different optimal policies when
 463 system dynamics are known perfectly (Figure 1B).




Figure 4: A. Model set representing structural uncertainty in the per-capita growth rate. B. Optimal state-dependent removal actions assuming perfectly known population dynamics, solved with stochastic dynamic programming. In both panels, colors correspond to different values of m in the model set, and the red dashed line corresponds to the true model dynamics used for simulation in the passive adaptive management approach.

464 Value of Information

465 For a decision maker, uncertainty may not be important to resolve; learning may not be
 466 inherently valuable if it does not result in improved management outcomes. A Value of
 467 Information (VOI) analysis *a priori* quantifies the importance of reducing uncertainty, or the
 468 expected improvement on management objectives if uncertainty were resolved (?). If the VOI is
 469 high, an adaptive management approach is typically justified.

470 The expected value of perfect information (EVPI) is measured as the difference between 1) PI ,
 471 or the expected value once uncertainty has been resolved because the optimal action is chosen

after knowing which model best describes the system, and 2) NL , or the expected value in the face of uncertainty because the optimal action is chosen with no learning about the system (?). The difference between these terms, $EVPI = PI - NL$, is the expected improvement in management performance due to acquiring perfect information.

To calculate EVPI for a sequential decision problem, we first use SDP to calculate the optimal state-dependent policy with assumed system dynamics, $\pi^*(s)_k$, for each model k in the model set K , given each transition matrix, P_k . We also calculate the bet-hedging strategy, or the optimal policy given uncertain dynamics, $\pi^*(s)_{NL}$. This bet-hedging strategy is calculated with the transition matrix P_{NL} , representing the average of the transitions of each model, P_k , weighted by each model belief, b : $P_{NL}(s_{t+1}|s_t, a_t) = \sum_{k \in K} b(k)P_k(s_{t+1}|s_t, a_t)$.

Using the policies calculated with known dynamics, $\pi^*(s)_k$, and unknown dynamics, $\pi^*(s)_{NL}$, we find the EVPI by first simulating i replicated population trajectories assuming each model k , applying both policies with known and unknown dynamics, and calculating the value at time t over time horizon T . The value associated with applying the policy with known dynamics is $V_{k,t,i}(s_t, a_t|a_t = \pi^*(s)_k)$, and the value associated with applying unknown dynamics is $V_{k,t,i}(s_t, a_t|a_t = \pi^*(s)_{NL})$. The EVPI is therefore $E_{s,k,t,i}[V(s_t, a_t|a_t = \pi^*(s)_k)] - E_{s,k,t,i}[V(s_t, a_t|a_t = \pi^*(s)_{NL})]$.

In the context of coypru regulation, the EVPI represents the expected improvement on the management objective if uncertainty about the per-capita growth rate was reduced. We calculate the EVPI for a model set, K , where $m \in \{0.4, 0.7, 1, 1.3, 1.7\}$ for $I = 200$ iterations over a $T = 40$ time horizon. We assume equal belief in each model. The expected value once uncertainty has been resolved, PI , is -1730.4, the expected value with no learning, NL , is -1787.8, and the EVPI is $PI - NL = 57.4$. It is therefore valuable to learn about the population dynamics and per-capita growth rate, as this learning is expected to improve performance on the management objective.

Value of information analyses are typically applied for static, one-time conservation decisions,

rather than sequential problems, often in the context of understanding when to act and when to delay action until after further research (?). To better understand the value of information in the context of adaptive management, future work should investigate quantifying VOI by iteratively resolving uncertainty, requiring evaluating the gain of implementing an adaptive policy rather than a single action (?).

Passive adaptive management

Passive adaptive management involves learning from the outcomes of previous management actions to update knowledge about system dynamics and improve management performance over time. Decisions at each time step are optimized assuming that current knowledge of the system will not change into the future. Learning therefore occurs by monitoring the effect of an implemented action, rather than during the optimization procedure (?).

Solving a passive adaptive management problem includes an implementation step and an updating step at every time point t , based on the current belief in model k , $b_t(k)$. The MDP must be solved at every time step because the transition matrices change as the belief changes. During the implementation step, the optimal action at time t is selected using a weighted averaging approach, where the transition probabilities are averaged across all models with the weights given by the current belief, $b_t(k)$ (?). This can be solved using stochastic dynamic programming:

$$V^*(s_t|b_t) = \max_{a \in A} \left\{ \sum_{k \in K} b_t(k) r(s_t, a_t) + \gamma \sum_{s_{t+1}} \left\{ \sum_{k \in K} b_t(k) P_k(s_{t+1}|s_t, a_t) \times V^*(s_{t+1}|b_t) \right\} \right\}$$

During the updating step, the resulting state, s_{t+1} , after taking action a_t is observed, and belief, b_{t+1} , is updated according to Bayes' rule:

$$b_{t+1}(k|s_t, a_t, s_{t+1}, b_t) = \frac{P_k(s_{t+1}|s_t, a_t)b_t(k)}{\sum_{k \in K} P_k(s_{t+1}|s_t, a_t)b_t(k)}$$

516 For coyapu control decisions, we solve the passive adaptive management problem by updating
 517 the belief about the coyapu population dynamics to improve management performance over
 518 time. We begin with an equal belief, $b_{t=1}(k) = 1/5$, across the five models in the model set
 519 (Figure 1A). We iterate between 1) calculating the optimal action given the current belief state, 2)
 520 implementing the optimal action and simulating the next state, s_{t+1} , given the true population
 521 dynamics, and 3) updating the belief state based observing the next state given action a_t . Here
 522 we assume the state is perfectly observed. We simulate the application of passive adaptive
 523 management over 30 time steps. The belief in each model changes over time (Figure 2A),
 524 resulting in the highest belief for the two models closest to the true population dynamics $m = 1$
 525 and $m = 1.3$ (Figure 1A). Through updating the belief state over time and implementing the
 526 optimal action given the current belief state, the management performance improves over time
 527 (Figure 2B).




Figure 5: *A.* Belief state for each model k at time t . Colors correspond to different values of m , representing structural uncertainty in the per-capita growth rate. *B.* Reward, $r(s_t, a_t)$, at time t in passive adaptive management problem.

Active adaptive management

In contrast to passive adaptive management, active adaptive management does not rely only on past experience and instead explicitly accounts for future learning opportunities to maximize the chance of achieving the objective. Instead of an optimization at every time step as in passive adaptive management, the optimization occurs once and requires that the trajectory of belief b_t to b_{t+1} and state transitions be calculated for all time steps (?). The optimal value function V^* that characterizes the performance of a policy is a function of both the state of the system s_t and belief over the models b_t :

$$V^*(s_t, b_t) = \max_{a \in A} \left\{ \sum_{k \in K} b_t(k) r(s_t, a_t) + \gamma \sum_{s_{t+1}} \left\{ \sum_{k \in K} b_t(k) P_k(s_{t+1} | s_t, a_t) \times V^*(s_{t+1}, b_t) \right\} \right\}$$

Active adaptive management with model uncertainty can be solved with many methods, including those developed for dealing with partial observability (?, ?). Here, a POMDP is characterized by a tuple $\langle X, A, O, P, Z, r, \gamma \rangle$, where $X = S \times K$ represents the factored state space with S denoting the possible conditions of the system and K denoting the model set. The unknown model is an unobserved state variable that must be inferred through observation. We can solve the POMDP with the Successive Approximations of the Reachable Space under Optimal Policies (SARSOP) algorithm (?). With this algorithm, the value function V^* associated with the optimal policy π^* can be approximated by a convex, piecewise-linear function $V(b) = \max_{\alpha \in \Gamma} (\alpha \cdot b)$, where Γ is a finite set of α -vectors and b is the discrete vector representation of a belief (?).

For coypru control decisions, we solve the active adaptive management problem by augmenting the state space with the belief in two models where $m \in \{0.7, 1.3\}$. We choose a smaller model set than in the passive adaptive management problem due to the computational cost associated with exploring the belief and state spaces over the entire time period (i.e., “curse of dimensionality”). We assume a uniform belief across models and assume perfect observation of the state (i.e., $o_t = s_t$).

We also use the coypru example to demonstrate the “dual control problem” (?). A management action altering a poorly understood system can produce two types of benefits: 1) short-term payoffs and 2) learning about the system so as to improve payoffs in the long-term. There is a trade-off between these two benefits, as large disturbances allow for rapid learning yet may incur high costs; the challenge is to optimally balance these benefits. Here we explore the dual

control problem by selecting two different discount factors, $\gamma \in \{0.95, 0.75\}$. The higher of which represents a decision maker who values long-term rewards more highly, and the lower of which represents a decision maker who values short-term rewards more highly.

The active adaptive management solution describes the optimal action (a_t , removal effort), given the state (s_t , coypu abundance), initial belief ($b_{t=1}$, uniform across models), and discount factor (Figure 3). The solution differs depending on how future rewards are discounted in the optimization. At low coypu abundance, the optimal removal effort increases as the discount factor increases. With a higher discount factor, future system states are valued more highly; even though a higher removal effort is costly today, this higher removal effort will accelerate learning to yield higher rewards in the long term (Figure 3A). In contrast, with a lower discount factor, active learning becomes less useful, as the price paid for not following the passive-adaptive (optimal right now) solution exceeds the information gained from learning (Figure 3B).




Figure 6: A. State-dependent active adaptive management solution, given an initial uniform belief across models and a higher discount factor representing a decision maker who values long-term payoffs more highly. B. State-dependent active adaptive management solution, given an initial uniform belief across models and a lower discount factor representing a decision maker who values short-term payoffs more highly.

569 On a single species ; check out Conroy's book, there is a nice figure illustrating AM.

570 <https://onlinelibrary.wiley.com/doi/10.1002/ece3.71176>

571 **Conclusion (all)**

572 Faire une conclusion courte, voir les exemples de How-to. Rappeler les main messages, how is

573 SDM/SDP relevant et qqsperspectives

574 **Acknowledgments**

575 **Data availability statement**

576 Data and code are available at <https://github.com/oliviergimenez/SDPaper>.

577 **References**

578 Baxter, P. W. J., and H. P. Possingham. 2011. Optimizing search strategies for invasive pests:

579 Learn before you leap. *Journal of Applied Ecology* 48:86–95.

580 Bogich, T. L., A. M. Liebhold, and K. Shea. 2008. To sample or eradicate? A cost minimization

581 model for monitoring and managing an invasive species. *Journal of Applied Ecology*

582 45:1134–1142.

583 Bolam, F. C., M. J. Grainger, K. L. Mengersen, G. B. Stewart, W. J. Sutherland, M. C. Runge,

584 and P. J. McGowan. 2019. Using the value of information to improve conservation decision

585 making. *Biological Reviews* 94:629–647.

586 Chadès, I., S. Nicol, T. M. Rout, M. Péron, Y. Dujardin, J.-B. Pichancourt, A. Hastings, and C. E.

587 Hauser. 2017. Optimization methods to solve adaptive management problems. *Theoretical*

588 *Ecology* 10:1–20.

589 Conroy, M. J., and J. T. Peterson. 2013. Decision making in natural resource management: A

structured, adaptive approach. John Wiley & Sons.

Crowley, S. L., S. Hinchliffe, and R. A. McDonald. 2017. Invasive species management will benefit from social impact assessment. *Journal of Applied Ecology* 54:351–357.

Estévez, R. A., C. B. Anderson, J. C. Pizarro, and M. A. Burgman. 2015. Clarifying values, risk perceptions, and attitudes to resolve or avoid social conflicts in invasive species management. *Conservation Biology* 29:19–30.

Fowler, C. W. 1981. Density dependence as related to life history strategy. *Ecology* 62:602–610.

Haubrock, P. J., T. Everts, N. A. S. Abreo, J. Bojko, V. Deklerck, J. W. E. Dickey, A. C. S. Franco, E. García-Berthou, S. Katsanevakis, N. I. Kirichenko, S. Mammola, M. A. Nuñez, B. Parker, R. Scalera, I. Soto, D. Strubbe, A. S. Tarkan, L. Vilizzi, T. Adriaens, P. Balzani, D. Blońska, E. Briski, R. Brys, A. L. Burgess, J. E. Byers, C. Cano-Barbacil, G. Castaldelli, J. T. A. Dick, V. Dominguez Almela, R. D. Dimarco, M. Florencio, A. Kouba, M. Kourantidou, I. Kurtul, I. Martín-Forés, O. Morissette, J. D. Olden, B. E. Soares, J. Truszkowski, H. Verreycken, M. Kenis, R. Sousa, and J. R. Britton. 2025. The impacts of biological invasions. *Biological Reviews* n/a.

Hauser, C. E., and H. P. Possingham. 2008. Experimental or precautionary? Adaptive management over a range of time horizons. *Journal of Applied Ecology* 45:72–81.

Holling, C. S. 1978. Adaptive environmental assessment and management. John Wiley & Sons.

Hyder, A., B. Leung, and Z. Miao. 2008. Ecology and Society: Integrating Data, Biology, and Decision Models for Invasive Species Management: Application to Leafy Spurge (*Euphorbia Esula*). *Ecology and Society* 13.

Hyytiäinen, K., M. Lehtiniemi, J. K. Niemi, and K. Tikka. 2013. An optimization framework for addressing aquatic invasive species. *Ecological Economics* 91:69–79.

Kurniawati, H., D. Hsu, W. S. Lee, and others. 2008. Sarsop: Efficient point-based pomdp planning by approximating optimally reachable belief spaces. *Robotics: Science and systems*. Zurich, Switzerland.

Leung, B., D. M. Lodge, D. Finnoff, J. F. Shogren, M. A. Lewis, and G. Lamberti. 2002. An ounce of prevention or a pound of cure: Bioeconomic risk analysis of invasive species. *Proceedings. Biological Sciences* 269:2407–2413.

Marescot, L., G. Chapron, I. Chadès, P. L. Fackler, C. Duchamp, E. Marboutin, and O. Gimenez. 2013. Complex decisions made simple: A primer on stochastic dynamic programming. *Methods in Ecology and Evolution* 4:872–884.

Pyšek, P., P. E. Hulme, D. Simberloff, S. Bacher, T. M. Blackburn, J. T. Carlton, W. Dawson, F. Essl, L. C. Foxcroft, P. Genovesi, J. M. Jeschke, I. Kühn, A. M. Liebhold, N. E. Mandrak, L. A. Meyerson, A. Pauchard, J. Pergl, H. E. Roy, H. Seebens, M. van Kleunen, M. Vilà, M. J. Wingfield, and D. M. Richardson. 2020. Scientists’ warning on invasive alien species. *Biological Reviews* 95:1511–1534.

Raiffa, H., and R. Schlaifer. 2000. *Applied statistical decision theory*. John Wiley & Sons.

Roy, H. E., A. Pauchard, P. J. Stoett, T. Renard Truong, L. A. Meyerson, S. Bacher, B. S. Galil, P. E. Hulme, T. Ikeda, S. Kavileveetttil, M. A. McGeoch, M. A. Nuñez, A. Ordonez, S. J. Rahlao, E. Schwindt, H. Seebens, A. W. Sheppard, V. Vandvik, A. Aleksanyan, M. Ansong, T. August, R. Blanchard, E. Brugnoli, J. K. Bukombe, B. Bwalya, C. Byun, M. Camacho-Cervantes, P. Cassey, M. L. Castillo, F. Courchamp, K. Dehnen-Schmutz, R. D. Zenni, C. Egawa, F. Essl, G. Fayvush, R. D. Fernandez, M. Fernandez, L. C. Foxcroft, P. Genovesi, Q. J. Groom, A. I. González, A. Helm, I. Herrera, A. J. Hiremath, P. L. Howard, C. Hui, M. Ikegami, E. Keskin, A. Koyama, S. Ksenofontov, B. Lenzner, T. Lipinskaya, J. L. Lockwood, D. C. Mangwa, A. F. Martinou, S. M. McDermott, C. L. Morales, J. Müllerová, N. A. Mungi, L. K. Munishi, H. Ojaveer, S. N. Pagad, N. P. K. T. S. Pallewatta, L. R. Peacock, E. Per, J. Pergl, C. Preda, P. Pyšek, R. K. Rai, A. Ricciardi, D. M. Richardson, S. Riley, B. J. Rono, E. Ryan-Colton, H. Saeedi, B. B. Shrestha, D. Simberloff, A. Tawake, E. Tricarico, S. Vanderhoeven, J. Vicente, M. Vilà, W. Wanzala, V. Werenkraut, O. L. F. Weyl, J. R. U. Wilson, R. O. Xavier, and S. R. Ziller. 2024. Curbing the major and growing threats from invasive alien species is urgent

and achievable. *Nature Ecology & Evolution* 8:1216–1223.

Runge, M. C., S. J. Converse, and J. E. Lyons. 2011. Which uncertainty? Using expert elicitation and expected value of information to design an adaptive program. *Biological Conservation* 144:1214–1223.

Runge, M. C., S. J. Converse, J. E. Lyons, and D. R. Smith (Eds.). 2020. *Structured Decision Making: Case Studies in Natural Resource Management*. Johns Hopkins University Press, Baltimore, Maryland.

Simberloff, D., J.-L. Martin, P. Genovesi, V. Maris, D. A. Wardle, J. Aronson, F. Courchamp, B. Galil, E. García-Berthou, M. Pascal, P. Pyšek, R. Sousa, E. Tabacchi, and M. Vilà. 2013. Impacts of biological invasions: What’s what and the way forward. *Trends in Ecology & Evolution* 28:58–66.

Thompson, B. K., J. D. Olden, and S. J. Converse. 2021. Mechanistic invasive species management models and their application in conservation. *Conservation Science and Practice* 3:e533.

Walters, C. J. 1986. *Adaptive management of renewable resources*. Macmillan Publishers Ltd.

Walters, C. J., and R. Hilborn. 1978. Ecological optimization and adaptive management. *Annual review of Ecology and Systematics* 9:157–188.

Williams, B. K., M. J. Eaton, and D. R. Breininger. 2011. Adaptive resource management and the value of information. *Ecological Modelling* 222:3429–3436.

Williams, B. K., J. D. Nichols, and M. J. Conroy. 2002. *Analysis and management of animal populations*. Academic Press, San Diego, California, USA.