

A primer on structured decision making for invasive species (using dynamic programming)

Olivier Gimenez^{1*} Abby Keller^{2*} Lucile Marescot^{3*} Cassie Speakman^{4*}

2026-02-27

¹ CEFE, Univ Montpellier, CNRS, EPHE, IRD, Montpellier, France ² address ³ address ⁴ address

* Corresponding author: olivier.gimenez@cefe.cnrs.fr

Invasive species (IAS) pose significant ecological, economic, and public health challenges.

Robust methods are needed for guiding effective management. Overall, we hope to demonstrate the flexibility and utility of structured decision making in invasive species management.

Target journal: Biological invasions? NeoBiota?.

Keywords: Adaptive management, Invasive species, Stochastic dynamic programming, Structured decision making

Introduction

Invasive species are a significant global issue, with wide-ranging impacts on ecosystems, economies, and public health (Pyšek et al. 2020, Roy et al. 2024).

Effective management of invasive species requires blabla for guiding control efforts and evaluating the success of eradication or regulation programs (Williams et al. 2002, Thompson et al. 2021).

Pitch: Structured decision-making (SDM) and adaptive management (AM) have been repeatedly recommended for the management of invasive species, yet they remain underused in practice. There are many reasons for this gap, one of which is the technical challenge of formally identifying and quantifying trade-offs among management objectives. Here, we contribute to ongoing efforts by illustrating the use of stochastic dynamic programming (SDP) and some of its variants — including active and passive adaptive management and partially observable Markov decision processes (POMDPs) — across different ecological contexts (a single population, a metapopulation, and two interacting populations).

Several papers have recommended or implemented structured decision-making (SDM) for managing invasive species. It's worth providing a short review of this literature.

In a recent review on impacts, the use of ML to guide management actions identified as a future research direction (see Appendix 3 in Haubrock et al. (2025)).

“Another recent research direction is the use of machine learning to aid the development of conservation management strategies. These methods involve models that integrate geospatial, climate, species distribution and other data to simulate the distribution of species and vegetation types under different land use and conservation planning scenarios. Reinforcement Learning algorithms can then be used to learn a management policy that attains a specified quantifiable objective (such as minimising the number of individuals from a non-native species) under

constraints such as cost or the amount of land that can be protected. Dietterich, Taleghan and Crowley (2013) demonstrated how Reinforcement Learning can be applied to managing a river system to control the expansion of a non-native plant species. Lapeyrolerie et al. (2022) demonstrated the potential of Reinforcement Learning for resource management problems, such as setting fishing quotas. Conservation area prioritisation through artificial intelligence (CAPTAIN) (Silvestro et al., 2022) is a general simulation framework that allows the user to define flexibly management objectives to be optimised. The unprecedented rate at which new technologies are emerging holds great promise for improving impact assessments related to biological invasions.”

Coypu (*Myocastor coypu*) as a case study throughout the paper. Code and tutorials available as appendices.

Add a shiny for each case study so that folks can play around?

A primer on structured decision making

We do not want to rewrite Lucile’s paper...

This is the expected discounted future value; If I take action a now in some state, this is my average (hence the sum) of tomorrow's best-possible rewards, scaled down by β because tomorrow's payoffs are worth less than


⊕

$$V(X_t) = \max_{A_t} \left[R_t(X_t, A_t) + \beta \sum_{X_{t+1}} \Pr(X_{t+1} | X_t, A_t) V(X_{t+1}) \right]$$

The immediate reward you get right now by taking action A in state X.

You choose the action A (among all available actions A) that maximizes the quantity inside the brackets.

The best (maximum) expected total reward you can collect from time t onward if you start in state X



53

54 Stochastic dynamic programming (Olivier)

55 Here we present a minimal stochastic dynamic programming (SDP) example. The aim is purely
 56 pedagogical: to show how states, actions, transition probabilities, rewards, and time horizon
 57 combine to produce an optimal decision rule.

58 We consider a stylized invasion management problem with three qualitative states describing
 59 the invasion level:

- 60 • Eradicated: no individuals detected (possible re-invasion risk)
- 61 • Contained: invasion present but at low abundance
- 62 • Established: invasion widespread with high impacts

63 At each time step, the manager chooses between two actions:

- Monitor: surveillance only (low cost, little effect)
- Control: active management campaign (higher cost, improves state)

State dynamics are described by action-dependent transition matrices

$$P_a(s, s') = \Pr(S_{t+1} = s' \mid S_t = s, A_t = a).$$

Transition matrices describe how the invasion state evolves probabilistically from one time step to the next under each management action. Each row corresponds to the current state and each column to the next state, with entries giving transition probabilities. We build a transition matrix for each action.

Under monitoring, the transition is:

		Erad.	Cont.	Estab.
$P_{\text{Monitor}} =$	Erad.	0.95	0.05	0.00
	Cont.	0.00	0.80	0.20
	Estab.	0.00	0.00	1.00

If the system is eradicated, it remains so with probability 0.95 but may revert to a contained state with probability 0.05, reflecting a small risk of re-invasion. When the invasion is contained, it persists with probability 0.80 but may deteriorate to an established state with probability 0.20, indicating possible population growth without intervention. Once established, the system remains established with probability 1, meaning monitoring alone cannot reverse the invasion.

In contrast, control shifts the system toward better states:

		Erad.	Cont.	Estab.
$P_{\text{Control}} =$	Erad.	0.98	0.02	0.00
	Cont.	0.20	0.75	0.05
	Estab.	0.05	0.75	0.20

From an eradicated state, the system remains eradicated with probability 0.98 and only rarely becomes contained (0.02). When contained, control can eradicate the population (0.20), maintain it at low levels (0.75), or fail and allow establishment (0.05). Finally, when the invasion is established, control improves the situation in most cases, leading to a contained state with probability 0.75, occasionally achieving eradication (0.05), but sometimes failing to reduce impacts (remaining established with probability 0.20).

Together, these transition matrices encode both ecological uncertainty and the imperfect effectiveness of management actions.

Next we define rewards. Rewards are defined as the negative of total costs, so that maximizing expected reward is equivalent to minimizing management and damage costs.

$$R(s, a) = -(\text{damage}(s) + \text{cost}(a)).$$

Damage depends on the invasion state (higher when the invasion is established), while management cost depends on the chosen action (higher for active control than monitoring). The reward matrix is constructed by combining these two components for all state–action pairs, yielding the immediate payoff used in the dynamic programming recursion.

		Monitor	Control
$R =$	Erad.	-1	-6
	Cont.	-5	-10
	Estab.	-15	-20

(Rows correspond to the current invasion state and columns to the management action.)

For example, when the invasion is eradicated, monitoring yields a reward of -1, reflecting only minor surveillance costs, whereas applying control in the same state yields -6 due to higher intervention costs. When the population is contained, rewards are -5 for monitoring and -10 for control, representing moderate damages combined with the respective management costs. Finally, when the invasion is established, rewards drop to -15 under monitoring and -20 under control, reflecting both high ecological impacts and, in the latter case, substantial intervention effort. Overall, this reward structure captures the short-term trade-off between ecological impacts and management effort that underlies the decision problem.

Solving the finite-horizon decision problem yields the optimal value function $V_t(s)$, which represents the expected cumulative reward (i.e., negative total cost) obtained when starting from state s at time t and following the optimal policy thereafter. Values are more negative when the invasion is established, reflecting higher expected damages. Moreover, values increase (become less negative) as the horizon shortens because fewer future costs can accumulate.

		$t = 1$	$t = 2$	$t = 3$	$t = 4$	$t = 5$	$t = 6$
$V_t(s) =$	$s = \text{Erad.}$	-7.60	-5.50	-3.69	-2.20	-1	0
	$s = \text{Cont.}$	-33.78	-27.43	-19.96	-12.0	-5	0
	$s = \text{Estab.}$	-49.26	-42.05	-34.47	-26.8	-15	0

The associated optimal policy $\pi_t(s)$ specifies the best action to take in each state s at each

time step t . Overall, the policy indicates that monitoring is typically optimal when the system is eradicated, whereas control is recommended when the invasion is established due to the high expected damages. In the contained state, the optimal decision varies with the time horizon, reflecting the trade-off between immediate intervention costs and the risk of future deterioration.

	$t = 1$	$t = 2$	$t = 3$	$t = 4$	$t = 5$
$\pi_t(s) =$					
$s = \text{Erad.}$	Monitor	Monitor	Monitor	Monitor	Monitor
$s = \text{Cont.}$	Control	Control	Monitor	Monitor	Monitor
$s = \text{Estab.}$	Control	Control	Control	Control	Monitor

This toy example highlights key features of SDP:

- **State dependence:** optimal decisions depend on invasion level.
- **Trade-off:** control is costly but reduces future risk.
- **Time horizon effects:** as the horizon shortens, actions may become more short-term oriented.
- **Uncertainty:** decisions account for probabilistic transitions, not deterministic outcomes.

Case studies with coypus

Single population (Olivier)

The objective of this example is to determine an optimal regulation policy for coypu based on their total abundance. We consider a stylized management problem in which a decision maker chooses, at each time step, the intensity of a control effort $u \in [0, 1]$. Low effort corresponds to minimal intervention, while high effort represents intensive removal. The management objective is to balance three competing components: ecological damages that increase with

population size (e.g. crop losses, dike weakening, nuisance), management costs that increase with effort, and an additional penalty if abundance exceeds a tolerable level.

In contrast to the previous example with discrete invasion states, the system is here described by a continuous state variable, the population size N . Because dynamic programming requires a finite state space, abundance is discretized onto a grid ranging from 0 to the carrying capacity K . In the implementation, we set $K = 2000$ individuals and use a grid step of 20 individuals. Control effort is also discretized, from 0 to 1 in increments of 0.05.

Population dynamics follow a density-dependent growth model with removals due to control,

$$N_{t+1} = N_t + rN_t \left(1 - \frac{N_t}{K}\right) - q(u)N_t,$$

where the first term describes logistic growth and $q(u)$ is the fraction of individuals removed as a function of effort. We use a somehow saturating removal function,

$$q(u) = 1 - \exp(-\alpha u),$$

with $\alpha = 1.5$, which implies diminishing returns of control effort (Figure 1). For example, an effort of $u = 0.5$ removes about 53% of the population, whereas maximum effort $u = 1$ removes about 78%, reflecting the fact that complete eradication is rarely achievable in a single step.

Rewards are defined as the negative of total costs, so that maximizing expected reward is equivalent to minimizing long-term damages and management expenses (Figure 2):

$$R(N, u) = -(\text{damage}(N) + \text{cost}(u) + \text{penalty}(N)).$$

Damages increase non-linearly with abundance, here using a simple quadratic form $\text{damage}(N) = 0.1N + 0.001N^2$, reflecting the common ecological pattern whereby impacts




Figure 1: Removal fraction $q(u)$ as a function of control effort u .

142 accelerate at high density. Management costs increase with effort and are assumed convex
 143 ($\text{cost}(u) = 50u + 200u^2$), capturing increasing marginal costs of intensive interventions. Finally,
 144 a penalty is applied when abundance exceeds a tolerable threshold $N_{\text{tol}} = 400$, representing
 145 situations where ecological or social impacts become unacceptable. The penalty increases
 146 linearly beyond this threshold according to

$$\text{penalty}(N) = \begin{cases} 0, & N \leq N_{\text{tol}}, \\ 2000 \frac{N - N_{\text{tol}}}{N_{\text{tol}}}, & N > N_{\text{tol}}. \end{cases}$$

147 This formulation implies that the penalty is zero below the threshold and then rises proportion-
 148 ally to the exceedance. For example, an abundance of $N = 800$ yields a penalty of 2000, thereby
 149 strongly incentivizing control once populations grow beyond acceptable levels.

150 The decision problem is solved using infinite-horizon discounted value iteration (with discount
 151 factor $\gamma = 0.99$). The resulting value function represents the expected long-term net benefit of




Figure 2: Immediate cost components as functions of abundance. Damage increases nonlinearly with population size, leading to rapidly rising total costs (damage + penalty). The dotted line marks the tolerable threshold ($N_{tol} = 400$) above which an additional penalty is applied in the decision model, although it is not shown here for clarity. Control cost is illustrated for an example effort level ($u = 0.6$).

starting from each abundance level when optimal decisions are applied thereafter. The optimal policy $u(N)$ typically exhibits a threshold-like pattern: low effort at low abundance where damages are limited, progressively stronger intervention as abundance increases, and intensive control once the population approaches or exceeds the tolerable level. This behavior reflects a general ecological intuition — management intensity should scale with both population size and risk — and illustrates how stochastic dynamic programming formalizes this principle within a transparent decision framework.

The decision problem is solved using infinite-horizon discounted value iteration (discount factor $\gamma = 0.99$). This yields (i) an optimal stationary policy, mapping each abundance class to a recommended control effort, and (ii) a value function (not shown) that summarizes the long-term expected net benefit when the policy is followed.

Because damages and the penalty increase sharply with abundance, the optimal rule is “threshold-like” (top panel in Figure 3). Effort is low at small N , then increases rapidly as N approaches the tolerable threshold $N_{\text{tol}} = 400$, and remains high for large N . Ecologically, this tells us that when the population is small, impacts are limited and additional control has diminishing returns, whereas once abundance enters an unacceptable zone, strong intervention is warranted to prevent sustained high impacts.

The bottom panel of Figure 3 illustrates the dynamical consequences of applying this rule to a trajectory starting from $N_0 = 800$. Control effort $u(t)$ adjusts immediately to the current abundance class, and abundance $N(t)$ responds through the logistic dynamics with removals. In this deterministic grid-based version, the system typically moves quickly toward a lower-abundance regime and then stabilizes around a band where the marginal benefit of additional control balances its marginal cost.




Figure 3: Optimal policy $u^*(N)$ (top) and an example trajectory under the optimal policy (bottom, faceted into abundance and effort), starting from $N_0 = 800$.

Two interacting species (Lucile)

Competition or predator-prey, or host-pathogen / disease thing. See SDP-species-interactions / in Olivier biblio repo.

Metapopulation (Olivier)

We extend the single-population coy-pu example to a spatially structured invasion in which several sites are connected through dispersal. The aim is to explore how management effort should vary through time when local demographic processes and spatial spread interact. We consider a network of several sites that can be interpreted as cities or management units. At each monthly time step, a decision maker chooses a global management action corresponding to a harvest intensity applied simultaneously across all sites. Local populations then grow according to density dependence, are partially removed by management, and may subsequently undergo stochastic extinction or colonization depending on local abundance and the occupancy of neighboring sites.

At time t , each site i is either occupied or empty, represented by the binary variable $X_i(t) \in \{0, 1\}$. When a site is empty its abundance is zero, whereas when it is occupied its population size follows discrete logistic growth with intrinsic rate r and carrying capacity K . Denoting abundance by $N_i(t)$, post-growth abundance is

$$N_i^{\text{grow}}(t+1) = N_i(t) + r N_i(t) \left(1 - \frac{N_i(t)}{K}\right).$$

Management then removes a fraction h_a of the population, so that

$$N_i(t+1) = N_i^{\text{grow}}(t+1) (1 - h_a).$$

193 This formulation captures the idea that control reduces numbers but does not eliminate under-
 194 lying demographic regulation.

195 Occupancy is updated stochastically based on local abundance and spatial context (Figure ??).

196 When a site is occupied, extinction probability decreases with abundance according to

$$p_{\text{ext}}(N_i) = \exp(-\alpha N_i),$$


197 so that small populations are more vulnerable to disappearance. When a site is empty, coloniza-
 198 tion probability increases with the number of occupied neighbors

$$k_i(t) = \sum_{j \neq i} A_{ij} X_j(t),$$

199 through

$$p_{\text{col}}(k_i) = 1 - (1 - \beta)^{k_i}.$$

200 Together, these rules link local demography with spatial dynamics and create feedbacks between
 201 sites.



203 Because the transition probabilities induced by these ecological rules are analytically intractable,
 204 we approximate them using repeated simulations. Starting from a given state and management
 205 action, we simulate the ecological dynamics many times and record the frequency of each
 206 possible next state. In practice, the estimated probability of transitioning from state s to state s'
 207 under action a is the proportion of simulations that lead to s' :

208 This Monte Carlo approach allows the decision model to incorporate nonlinear and stochastic
 209 ecological processes without requiring an explicit analytical expression for transition probabili-
 210 ties.

211 Management decisions are evaluated through a cost function that balances ecological damages
 212 and intervention effort (Figure 4). At each time step, total cost is defined as

$$C(s, a) = c_d \sum_{i=1}^n N_i(t) + c_a(a),$$

213 where c_d represents the damage cost per individual (here $c_d = 1$) and $c_a(a)$ is the management
 214 cost associated with action a . In this example, we consider three possible actions corresponding
 215 to no control, moderate control, and strong control, with costs

$$c_a(a) \in \{0, 5, 10\}.$$

216 The reward used in the dynamic programming recursion is defined as the negative of this cost,

$$R(s, a) = -C(s, a),$$

217 so that maximizing expected reward is equivalent to minimizing long-term total costs.

218 The resulting policy is illustrated in Figure 5. When few sites are occupied and populations
 219 remain small, limited intervention is often sufficient because damages are modest and extinction




Figure 4: Representation of the cost components on the log scale. Damage costs increase linearly with total abundance across sites, whereas control costs are fixed action-specific intercepts (none/moderate/strong). Stronger control increases immediate spending but can reduce future abundance and damages.

may occur naturally. As more sites become occupied or local abundances increase, stronger control becomes optimal in order to reduce future spread and prevent escalating impacts. Temporal variation in recommended actions arises because the finite planning horizon reduces the benefit of aggressive intervention late in the time period.

Comparing forward simulations under the optimal policy with simulations under constant strategies highlights the value of adaptive management (Figure 6). The optimal strategy typically reduces cumulative costs by applying effort strategically in time rather than maintaining a fixed level of control.

Spatial trajectories further illustrate how coordinated management alters invasion dynamics across the network, with reductions in abundance and occupancy emerging progressively across sites (Figure 7).




Figure 5: Optimal global action as a function of time and invasion level (measured as the number of occupied sites). Colors indicate the action recommended by the finite-horizon SDP (None/Moderate/Strong). This visualization highlights how stronger control becomes optimal as the invasion occupies more sites, and how recommendations may soften near the end of the planning horizon.




Figure 6: Cumulative costs over time under constant strategies (none, moderate, strong) versus the SDP-optimal policy. Costs combine damages proportional to total abundance and an action-specific control cost.




Figure 7: Site-level abundance dynamics under constant strategies and the SDP-optimal policy. Each panel corresponds to a site; trajectories highlight how spatial heterogeneity and network-driven recolonization can create uneven benefits of control across locations.

231 **Accounting for uncertainty**

232 **POMDP (Cassie)**

233 Illustrate POMDP, Olivier has some code extending the computer maintenance example from
234 above.

235 <https://besjournals.onlinelibrary.wiley.com/doi/10.1111/2041-210X.14374>

236 <https://besjournals.onlinelibrary.wiley.com/doi/10.1111/2041-210X.13692>

237 **Adaptive management (Abby)**

238 When is it useful. Start with value of information?

239 On a single species ; check out Conroy's book, there is a nice figure illustrating AM.

240 <https://onlinelibrary.wiley.com/doi/10.1002/ece3.71176>

241 **Discussion (all)**

242 **Main messages**

243 How is SDM relevant.

244 **Limitations**

245 Curse of dimensionality.

Perspectives

Acknowledgments

Data availability statement

Data and code are available at <https://github.com/oliviergimenez/SDPaper>.

References

- Haubrock, P. J., T. Everts, N. A. S. Abreo, J. Bojko, V. Deklerck, J. W. E. Dickey, A. C. S. Franco, E. García-Berthou, S. Katsanevakis, N. I. Kirichenko, S. Mammola, M. A. Nuñez, B. Parker, R. Scalera, I. Soto, D. Strubbe, A. S. Tarkan, L. Vilizzi, T. Adriaens, P. Balzani, D. Blońska, E. Briski, R. Brys, A. L. Burgess, J. E. Byers, C. Cano-Barbacil, G. Castaldelli, J. T. A. Dick, V. Dominguez Almela, R. D. Dimarco, M. Florencio, A. Kouba, M. Kourantidou, I. Kurtul, I. Martín-Forés, O. Morissette, J. D. Olden, B. E. Soares, J. Truszkowski, H. Verreycken, M. Kenis, R. Sousa, and J. R. Britton. 2025. The impacts of biological invasions. *Biological Reviews* n/a.
- Pyšek, P., P. E. Hulme, D. Simberloff, S. Bacher, T. M. Blackburn, J. T. Carlton, W. Dawson, F. Essl, L. C. Foxcroft, P. Genovesi, J. M. Jeschke, I. Kühn, A. M. Liebhold, N. E. Mandrak, L. A. Meyerson, A. Pauchard, J. Pergl, H. E. Roy, H. Seebens, M. van Kleunen, M. Vilà, M. J. Wingfield, and D. M. Richardson. 2020. Scientists' warning on invasive alien species. *Biological Reviews* 95:1511–1534.
- Roy, H. E., A. Pauchard, P. J. Stoett, T. Renard Truong, L. A. Meyerson, S. Bacher, B. S. Galil, P. E. Hulme, T. Ikeda, S. Kavileveettil, M. A. McGeoch, M. A. Nuñez, A. Ordonez, S. J. Rahlao, E. Schwindt, H. Seebens, A. W. Sheppard, V. Vandvik, A. Aleksanyan, M. Ansong, T. August, R. Blanchard, E. Brugnoli, J. K. Bukombe, B. Bwalya, C. Byun, M. Camacho-Cervantes, P.

Cassey, M. L. Castillo, F. Courchamp, K. Dehnen-Schmutz, R. D. Zenni, C. Egawa, F. Essl,
 G. Fayvush, R. D. Fernandez, M. Fernandez, L. C. Foxcroft, P. Genovesi, Q. J. Groom, A. I.
 González, A. Helm, I. Herrera, A. J. Hiremath, P. L. Howard, C. Hui, M. Ikegami, E. Keskin,
 A. Koyama, S. Ksenofontov, B. Lenzner, T. Lipinskaya, J. L. Lockwood, D. C. Mangwa, A.
 F. Martinou, S. M. McDermott, C. L. Morales, J. Müllerová, N. A. Mungi, L. K. Munishi,
 H. Ojaveer, S. N. Pagad, N. P. K. T. S. Pallewatta, L. R. Peacock, E. Per, J. Pergl, C. Preda,
 P. Pyšek, R. K. Rai, A. Ricciardi, D. M. Richardson, S. Riley, B. J. Rono, E. Ryan-Colton, H.
 Saeedi, B. B. Shrestha, D. Simberloff, A. Tawake, E. Tricarico, S. Vanderhoeven, J. Vicente,
 M. Vilà, W. Wanzala, V. Werenkraut, O. L. F. Weyl, J. R. U. Wilson, R. O. Xavier, and S. R.
 Ziller. 2024. Curbing the major and growing threats from invasive alien species is urgent
 and achievable. *Nature Ecology & Evolution* 8:1216–1223.

Thompson, B. K., J. D. Olden, and S. J. Converse. 2021. Mechanistic invasive species man-
 agement models and their application in conservation. *Conservation Science and Practice*
 3:e533.

Williams, B. K., J. D. Nichols, and M. J. Conroy. 2002. Analysis and management of animal
 populations. Academic Press, San Diego, California, USA.